


SAVIO Cluster

Set up new accounts

1. Create portal accounts at MyBRC User Portal (<https://mybrc.brc.berkeley.edu/>)
2. After following the on-screen instructions in the portal and registering and/or logging in, you should first review and sign the cluster User Access Agreement Form on the Home ("Welcome") page (if they haven't already done so from within the portal previously) by clicking on the "Review" button.
3. Click on the "Join" button to request to join project, you can search “**ic_chem242**” under project name to find our class project. Click on "join".
4. We will approve your request. And for new users there will be a processing wait time of a few days or so between the time they join their first project via the MyBRC User Portal and the time they receive the confirmation email pointing them to instructions on how they can access their new BRC user account.
5. Set up one-time password authentication by following the instructions here: <https://docs-research-it.berkeley.edu/services/high-performance-computing/user-guide/setting-otp/>

Use Jupyter Notebook on Savio


1. Go to: <https://ood.brc.berkeley.edu/>
2. You will see the following prompts:



The screenshot shows a 'Sign in' dialog box with a light blue background. At the top, it says 'Sign in' followed by the URL 'https://ood.brc.berkeley.edu'. Below this are two input fields: 'Username' with the text 'ericwangyz' and 'Password' with masked characters '.....'. At the bottom right are two buttons: 'Cancel' and 'Sign In'.

Enter your BRC username and password to login. The password is your PIN followed by the 6-digit OTP. For details, please refer to: <https://docs-research-it.berkeley.edu/services/high-performance-computing/user-guide/logging-brc-clusters/>

3. After successful login, you will see the following page. Click the "Interactive Apps" pulldown and choose "Jupyter Server – compute via Slurm using Savio partitions"



4. Request a compute node by filling in the information as follows. You can change the "Name of the Job" and "Wall Clock Time" at your preference. But please don't waste computation time. Click "Launch" and wait for the session to start. If your session is queued for a long time, you can delete the waiting session and try again with smaller "Number of CPU Cores per Node" (2 or 4 is a reasonable choice).

Jupyter Server - compute via Slurm using Savio partitions

version: 2021.pumpkin-56-geb5e745

This app will launch a [Jupyter](#) server session on the Berkeley Research Computing (BRC) Savio cluster.

Name of the Job

chem242

☐ Use JupyterLab instead of Jupyter Notebook?

SLURM Partition

savio2_1080ti

The SLURM Partition in which you want to launch this Jupyter session. For code that is not parallelized (or using only a small number of CPU cores), the HTC partitions are a good choice.

SLURM Account/Project Name

ic_chem242

The SLURM account (i.e., the value of the -A or --account flag used when submitting a SLURM job).

SLURM QoS Name

savio_normal

The QoS you want run under. Most FCA users choose savio_normal. Condo users can either use their condo QoS or the low-priority QoS.

Number of Compute Nodes

1

The number of nodes you want for your Jupyter session.

Number of CPU Cores per Node

8

The number of CPU cores you want per node for your Jupyter session.

Number and Type of GPUs

gpu:1

The GRES value in the form 'gpu:x' or 'gpu:type:x' where 'x' is the number of GPUs and (optionally) 'type' is the type of GPU. Remember to specify two CPUs for each GPU in the 'Number of CPU Cores per Node' field.

Wall Clock Time

1

The maximum number of hours your Jupyter session will run for. To save FCA credits or free up resources for your condo group members, you should delete your session when you are done.

Email Address (Optional)

Enter your email address if you would like to receive an email when the session starts. Leave blank for no email.

Launch

5. You should see this page when the node is ready. Click on "Connect to Jupyter" to start working with Jupyter Notebook. If you finish earlier than the time you requested, you can click on Delete to free up the node for other people.

Jupyter Server - compute via Slurm using Savio partitions (17818297)

1 node | 8 cores | Running

Host: [>_n0226.savio2](#)

 Delete

Created at: 2024-03-30 22:04:52 PDT

Time Remaining: 58 minutes

Session ID: [02b5c192-2c7e-4758-8085-44e08d13b0d7](#)

 [Connect to Jupyter](#)